

OD - Distributed Graph Lab: Team OD-L2-11-G

MEYSAM ZAMANI FOROOSHANI - BOYAN ZHANG

meysam.zamani@est.fib.upc.edu - boyan.zhang@est.fib.upc.edu

Exercise 1: Getting familiar with GraphX's Pregel API

In order to make easier the understanding of the superstep explanations we will define the relation between some concepts and how they are represented in the following images:

1. Vertices:

- I The vertices are represented by the circles.
- II They are accompanied by a label that identifies them.
- III The circle color shows the current vertex state: Green ones are active and red ones are inactive.
- IV Inside the circle, we can find two numbers: the number with parentheses is the initial value and the number without parentheses is the current value for the vertex.

2. Edges:

- I The edges are represented with the grey thick arrows and shows the connections between the vertices.

3. Messages:

- I The messages sent between the vertices are represented with black thin arrows.
- II On the left figures, we show the messages that the vertices receive from the previous superstep or as init messages.
- III On the right figures, we show the messages that the vertices send to another one.
- IV The messages are labeled with a number that represents the message content.


Finally, we have two figures by each superstep. The left one is the initial one and the right one is the last one inside a superstep.

Superstep 0

In the first superstep, all the vertices are active because all receive the initial message.

- **Apply:** All vertices receive an initial message MAX_VALUE. Therefore for all vertices in the apply phase will just apply (return) their current values.
- **Scatter:** The scatter phase receives a triplet containing the source, destination and edge. The source vertex value is compared to the destination vertex value: a message is sent only if the source vertex value is greater than the destination vertex value. Only vertex 1 has any outgoing edge that satisfies this condition ($9 > 1$) and therefore is the only vertex that sends a message, namely 9 to vertex 2.
- **Gather:** All nodes receive just one message and so the gather phase is not used.

FIGURA 1 – SuperStep0

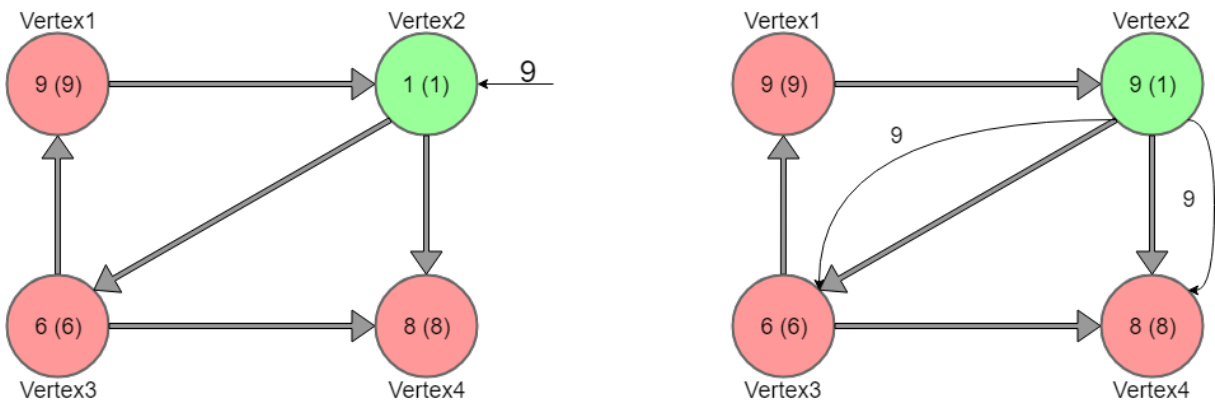


Superstep 1

On the second superstep, Vertex2 is the only active vertex.

- **Apply:** Vertex 2 evaluates the message by comparing against the current value. Given that the received value (9) is bigger than the current one (1), it updates its value to 9.
- **Scatter:** Vertex 2 has outgoing edges that satisfies the condition (see description in S0): $9 > 8$ and $9 > 6$. Therefore it sends messages to vertex 3 and 4 with the message 9.
- **Gather:** Vertex 2 receives just one message and so the gather phase is not used.

FIGURA 2 – SuperStep1




Superstep 2

On the third superstep, Vertex3 and Vertex4 are the active vertices.

- **Apply:** Vertex 3 and 4 evaluates their messages by comparing against the current value. Given that the received value (9) is bigger than the current ones (6 and 8), they update their values to 9.
- **Scatter:** Vertex 3 has outgoing edges but since their value is not greater than 9 (they are already 9), sends no message. Vertex 4 does not have any outgoing edges and finally the algorithm finishes.
- **Gather:** Vertex 3 and 4 receives just one message and so the gather phase is not used.

FIGURA 3 – SuperStep2



Exercise 2: Computing shortest paths using Prege

In order to clarify the code implemented in the three methods, we will describe it in the following sections.

- **VProg:** The apply method, in the VProg class, compares the current path cost and the received one in the message and select the smallest. In the first superstep, the value of the message always is the current path cost.
- **sendMsg:** The apply method, in the sendMsg class, checks if from the current vertex exists a better path to arrive at some neighbor vertex. So, if in this method the source vertex has an infinite number (it doesn't have any path to it) or the new possible path is worse than the current one, the vertex doesn't send the message. If instead, a new better path is possible, the vertex sends a message with the cost of this new path.
- **merge:** The apply method, in the merge class, takes two messages and select the message with the best path proposed (with the smallest cost).

Note: results are sorted by adding the following sortBy:

```
ops.pregel(Integer.MAX_VALUE,
    Integer.MAX_VALUE,
    EdgeDirection.Out(),
    new VProg(),
    new sendMsg(),
    new merge(),
    ClassTag$.MODULE$.apply(Integer.class)) Graph
    .vertices() VertexRDD
    .toJavaRDD() JavaRDD
    .sortBy(v1 -> labels.get(((Tuple2) v1)._1), ascending: true, numPartitions: 1)
    .foreach(v -> {
        Tuple2<Object, Integer> vertex = (Tuple2<Object, Integer>) v;
        System.out.println("Minimum cost to get from " + labels.get(11) + " to " + labels.get(vertex._1) + " is " + vertex._2);
    });
}
```

And finally we obtain the expected output like as the one written in the statement of the assignment.

Exercise 3: Extending shortest paths computation

In general terms, the algorithm is the same than in the previous exercise but saving the nodes that compose the path and not only the cost of the path.

Note: results are sorted by adding the following sortBy:

```
ops.pregel(new Tuple2("", Integer.MAX_VALUE),
    Integer.MAX_VALUE,
    EdgeDirection.Out(),
    new Exercise_3.VProg(),
    new Exercise_3.sendMsg(),
    new Exercise_3.merge(),
    ClassTag$.MODULE$.apply(Tuple2.class)) Graph
    .vertices() VertexRDD
    .toJavaRDD() JavaRDD
    .sortBy(v1 -> ((Tuple2) ((Tuple2) v1)._2)._1, ascending: true, numPartitions: 1)
    .foreach(v -> {
        Tuple2<Object, Tuple2> vertex = (Tuple2<Object, Tuple2>) v;
        System.out.println("Minimum cost to get from " + labels.get(11) + " to " + labels.get(vertex._1) +
    });
}
```

And finally we obtain the expected output like as the one written in the statement of the assignment except that our result is sorted by the cost number as well.

Exercise 4: Spark Graph Frames

The solution is very similar to the warmup exercise and The only different part from the warmup exercise aside from the pagerank function is the reading part. In order to make the code clean and simple we used the map function so that it applies the function spaceForBeauty to the edges and vertices.

You can find the screenshot of the result hereunder:

src	dst
36359329835505530	6843358505416683693
168437400931144903	961421098734626813
168437400931144903	1367968407401217879
168437400931144903	227043766454777682
168437400931144903	2381426201672413470
168437400931144903	3694025250309277803
168437400931144903	4289177408495534056
168437400931144903	4876378461651498574
168437400931144903	4898810637927996530
168437400931144903	5397795112799724806
168437400931144903	5513532497162072006
168437400931144903	5513532497446264506
168437400931144903	5781525238152196030
168437400931144903	6033170360494767837
168437400931144903	7647287135606994048
168437400931144903	8057641240036215591
168437400931144903	8069908956126960798
168437400931144903	8405829285832207680
179898883283219993	4315052278831489964
301281802531867654	8830299306937918434

id	name
6598434222544540151	Adelaide Hanscom ...
7814958205460279317	David Dodge (nove...
3858831448322232257	Howard League For...
1778261942684788432	Chelsea Quinn Yarb...
4201849915685975228	Dominick Montiglio
7447657998247897725	Joseph W. Estabrook
7816201114341688960	Spencer J. Palmer
290343697355110541	Aquatic Park (Ber...
1748906252821810168	Johan Galtung
3683742251178101007	Henry Howard, 7th...
8277480409175395091	Carl Dennis
8847001713698823343	Robert Axelrod
4011285656061867661	National Parliame...
1547980201223606474	Candice Nelson
5513532495900171287	Cartoon
1629895207124579982	Be Our Guest
7866692046732241596	Music of Zimbabwe
4898810267628502395	Scotland
6737425368859109805	Edward Clive, 1st...
1765549826147530806	Google China

id	name	pagerank
8830299306937918434	University of Cal...	3119.049280209878
1746517089350976281	Berkeley, California	1574.6334151225258
8262690695090170653	Uc berkeley	384.3346910461765
7097126743572404313	Berkeley Software...	214.47819575445723
8494280508059481751	Lawrence Berkeley...	194.84457483752416
1735121673437871410	George Berkeley	193.04164341777914
6990487747244935452	Busby Berkeley	110.93438197917992
1164897641584173425	Berkeley Hills	108.01090828891661
5820259228361337957	Xander Berkeley	70.85043533114853
6033170360494767837	Berkeley County, ...	67.72143888865237

only showing top 10 rows

Process finished with exit code 0